

Latency

- * Average latency: ~427 ms per image
- * Device: Apple MacBook Air (M2), PyTorch with Apple MPS acceleration

Cost per image

- * On-device deployment: Approximately \$0 per image, since inference runs locally and does not require cloud infrastructure.
- * Hypothetical cloud deployment: On the order of \$0.02–\$0.10 per 1,000 images, depending on the cloud provider, instance type, and utilization.